

```
In [2]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

```
In [2]: X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
```

```
In [3]: X.shape, y.shape
```

```
Out[3]: ((50000, 15), (50000,))
```

```
In [4]: #splitting dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

```
In [8]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [5]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[5]: ((37500, 15), (37500,), (12500, 15), (12500,))
```

SGD classifier

```
In [79]: clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
        clf
```

```
Out[79]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                       early_stopping=False, epsilon=0.1, eta0=0.0001,
                       fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                       loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                       penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                       tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [80]: clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.03 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.04 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.06 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.07 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.07 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.08 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.10 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.12 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.13 seconds.
Convergence after 10 epochs took 0.13 seconds
```

```
Out[80]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                        early_stopping=False, epsilon=0.1, eta0=0.0001,
                        fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                        loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                        penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                        tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [81]: clf.coef_, clf.coef_.shape, clf.intercept_
```

```
Out[81]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                  0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                  0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
         (1, 15),
         array([ -0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

```
In [3]: X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
```

```
In [4]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

Initialize weights

```
In [5]: def initialize_weights(dim):
```

```
    w=np.zeros_like(dim)
    b=0
    return w,b
```

```
In [6]: dim=X_train[0]
w,b = initialize_weights(dim)
print('w =',(w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function-1

```
In [7]: dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

Out[7]: True

Compute sigmoid

```
In [8]: def sigmoid(z):
sig=1/(1+np.exp(-z))
return sig
```

Grader function - 2

```
In [9]: def grader_sigmoid(z):
val=sigmoid(z)
assert(val==0.8807970779778823)
return True
grader_sigmoid(2)
```

Out[9]: True

Compute loss

```
In [10]: def logloss(y_true,y_pred):
n=len(y_true)
sum=0
for i in range(n):
    sum +=(y_true[i] * np.log10(y_pred[i])) + ((1 - y_true[i]) * np.log10(1 - y_pred[i]))
loss= -1*(1/n)*sum
return loss
```

Grader function - 3

```
In [11]: def grader_logloss(true,pred):
    loss=logloss(true,pred)
    assert(loss==0.07644900402910389)
    return True
true=[1,1,0,1,0]
pred=[0.9,0.8,0.1,0.8,0.2]
grader_logloss(true,pred)
```

Out[11]: True

Compute gradient w.r.to 'w'

```
In [12]: def gradient_dw(x,y,w,b,alpha,N):
    z=np.dot(w,x)+b
    dw=x*(y-sigmoid(z)-(alpha/N)*w)
    return dw
```

Grader function - 4

```
In [13]: def grader_dw(x,y,w,b,alpha,N):
    grad_dw=gradient_dw(x,y,w,b,alpha,N)
    assert(np.sum(grad_dw)==2.613689585)
    return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
    -2.81434437, -0.86771071, -0.04073287, 0.84827878, 1.99451725,
    3.67152472, 0.01451875, 2.01062888, 0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[13]: True

Compute gradient w.r.to 'b'

```
In [14]: def gradient_db(x,y,w,b):  
         z=np.dot(w,x)+b  
         db = y - sigmoid(z)  
         return db
```

Grader function - 5

```
In [15]: def grader_db(x,y,w,b):  
         grad_db=gradient_db(x,y,w,b)  
         assert(grad_db==-0.5)  
         return True  
         grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,  
                          -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,  
                          3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])  
         grad_y=0  
         grad_w,grad_b=initialize_weights(grad_x)  
         alpha=0.0001  
         N=len(X_train)  
         grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[15]: True

Implementing logistic regression

```
In [127]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
    train_loss=[]
    test_loss=[]
    dim=X_train[0]

    w,b = initialize_weights(dim)

    for i in range(epochs):

        for j in range(N):

            dw=gradient_dw(X_train[j],y_train[j],w,b,alpha,N)
            db=gradient_db(X_train[j],y_train[j],w,b)
            w = w + eta0 * dw
            b = b + eta0 * db

        train_pred=[sigmoid(np.dot(w, X_train[p]) + b) for p in range(len(X_train))]
        test_pred=[sigmoid(np.dot(w, X_test[q]) + b) for q in range(len(X_test))]

        loss1 = logloss(y_train, train_pred)
        train_loss.append(loss1)

        loss2 = logloss(y_test, test_pred)
        test_loss.append(loss2)

    return w,b,train_loss,test_loss
```

```
In [128]: alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=10
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```



```
In [130]: print("weights of custom implementation")  
w
```

weights of custom implementation

```
Out[130]: array([-0.42320237,  0.19097504, -0.14588903,  0.33813462, -0.21204108,  
                0.56528022, -0.44537758, -0.09169275,  0.21798654,  0.16980148,  
                0.19524869,  0.00226124, -0.0778474 ,  0.33881858,  0.02215503])
```

```
In [131]: print("intercept of custom implementation")  
b
```

intercept of custom implementation

```
Out[131]: -0.8505912823639118
```

Goal of assignment

```
In [132]: diff=np.subtract(w,clf.coef_)  
diff
```

```
Out[132]: array([[ 0.00016455,  0.00549939,  0.00270133, -0.00330945, -0.00385437,  
                0.00511443,  0.00704724,  0.00239538,  0.00871335, -0.01103978,  
               -0.00180321, -0.00195792,  0.0017563 ,  0.00029056, -0.00051218])
```

```
In [133]: k=[]  
          for list in diff:  
              for i in list:  
                  n = "{:e}".format(i)  
                  k.append(n)
```

```
In [134]: print("Difference of weights:")  
print(np.array(k))
```

Difference of weights:

```
['1.645475e-04' '5.499392e-03' '2.701329e-03' '-3.309448e-03'  
 '-3.854374e-03' '5.114433e-03' '7.047242e-03' '2.395376e-03'  
 '8.713346e-03' '-1.103978e-02' '-1.803215e-03' '-1.957920e-03'  
 '1.756299e-03' '2.905625e-04' '-5.121807e-04']
```

```
In [135]: print("Difference of intercepts:")  
np.subtract(b,clf.intercept_)
```

Difference of intercepts:


```
Out[135]: array([0.00254701])
```

Plot epoch number vs train , test loss

```
In [136]: from matplotlib import pyplot as plt
epoch = [i for i in range(1,11,1)]

plt.plot(epoch,train_loss , label='train_loss')
plt.plot(epoch,test_loss, label='test_loss')
plt.xlabel("epoch number")
plt.ylabel("log loss")
plt.legend()
plt.show
```

Out[136]: <function matplotlib.pyplot.show(*args, **kw)>



```
In [137]: def pred(w,b, X):  
    N = len(X)  
    predict = []  
    for i in range(N):  
        z=np.dot(w,X[i])+b  
        if sigmoid(z) >= 0.5:  
            predict.append(1)  
        else:  
            predict.append(0)  
    return np.array(predict)  
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))  
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))  
  
0.9553333333333334  
0.95288
```

```
In [ ]: !jupyter nbconvert --to html LRUsingSGD_solve.ipynb
```